

Building an Automated AI Travel Vlog

<https://gemini.google.com/app/f1989d0fd42db65c>

User prompt: Need to create a GenAI application using multimodal model with following feature 1. User can ingest a travel video. Bike ride car ride or simply roaming around to create a travel vlog. 2. Digital twin will put his overlay voice after analyzing the frames in video. location and detail about the place using web search. For this agentic application with tools at it disposal will be utilized. 3. The final video should have a great explanation with good background music as per the video frames. 4. To make it more realistic and exciting it should match with voice of person in frame for this AI model for audio will be utilized.


Response: This is a brilliant and highly ambitious architecture. Building a GenAI-powered **Automated AI Director** that takes raw footage and outputs a fully produced, context-aware travel vlog requires orchestrating multiple AI modalities (Vision, LLM Agents, Audio Cloning, and Video/Audio Editing).

To build this using an **Agentic Workflow**, you need a central controller (the Agent) managing specialized tools for video analysis, web research, voice cloning, and final synthesis.

Here is a comprehensive L400-level technical blueprint to design and implement this application.

1. High-Level Architecture & Workflow

The core of this system relies on decoupling **video comprehension** from **media generation**. Because LLMs cannot natively manipulate video timelines, the Agent acts as the brain, processing semantic data and generating an execution script for a programmatic video editing backend.



2. Component Break Down & Technical Stack

Step 1: Multimodal Ingestion & Frame Analysis

You need a model with a massive context window and strong native video understanding to handle long travel clips without losing context.

- **Core Model: Gemini 1.5 Pro.** Its 2-million token context window allows you to pass minutes of raw video directly as a native multimodal prompt, extracting temporal context without needing manual frame-splitting.
- **Alternative:** If using models like GPT-4o, you will need to sample frames (e.g., 1 frame per second) using OpenCV (cv2), encode them to Base64, and pass them as an array.
- **The Output:** The vision model must produce a **Structured Temporal Log** (JSON) detailing what is happening in the video at specific timestamps.

Step 2: The Orchestration Agent (LangGraph / CrewAI)

A simple linear pipeline won't work well here because a bike ride might pass an obscure monument that the vision model sees but can't identify precisely. The Agent needs to look at the visual description, formulate a search query, and enrich the context.

- **Framework: LangGraph** is highly recommended here over LangChain because video processing requires cyclical states (e.g., *Analyze Frame -> Search Web -> Refine Script -> Validate Timing*).
- **Agent Tools:**
 1. **Google Search API / Tavily API:** To fetch historical facts, local culture, and real-time data about the identified locations.

2. **Script Writer Engine:** A structured LLM call (using Pydantic output parsing) that outputs the final voiceover script mapped perfectly to timestamps.

Step 3: Voice Cloning & Alignment (Digital Twin Audio)

To match the voice of the person in the frame, you need a Few-Shot Text-to-Speech (TTS) engine.

- **Technology:** **ElevenLabs API** (for premium quality and instant voice cloning) or **Coqui XTTS v2 / Bark** (if you want an open-source, self-hosted option).
- **The Workflow:**
 1. The user uploads a 10-second reference audio clip of their voice during setup.
 2. The Agent feeds the generated timestamped script into the cloned voice model.
 3. **Crucial Step (Audio Pacing):** Travel videos have natural pauses. The TTS engine must generate separate audio files for each timestamp block so they can be precisely mapped to the video timeline without drifting.

Step 4: Contextual Background Music

Background music needs to match the vibe (e.g., high-energy electronic for a fast bike ride through a city, lo-fi acoustic for roaming a historic street).

- **Option A (Generative):** **Suno AI** or **Mubert API**. You pass the visual description as a prompt (e.g., *"Upbeat indie-folk acoustic guitar track for a scenic mountain drive"*).
- **Option B (Curated Selection):** A vector database (like Chroma or Pinecone) filled with categorized royalty-free tracks. The Agent generates a "music vibe description," converts it to an embedding, and finds the closest matching audio track from your library.

3. Implementation Plan: The Agentic Workflow (Python)

Here is a conceptual implementation pattern using Python, utilizing a LangGraph-like stateful approach for the Agent.

Phase 1: Temporal Video Analysis

```
import google.generativeai as genai

def analyze_travel_video(video_path):
    # Configure Gemini 1.5 Pro
    genai.configure(api_key="YOUR_GEMINI_API_KEY")
    model = genai.NotebookLM / genai.GenerativeModel('gemini-1.5-pro')

    # Upload video using the File API
    video_file = genai.upload_file(path=video_path)

    prompt = """
    Analyze this travel video and return a structured JSON array.
    Identify key locations, transitions, landmarks, and the overall 'vibe' of segments.
    Return schema:
    [
      {
        "start_time": "00:00",
        "end_time": "00:15",
        "visual_description": "Riding a motorcycle through a winding mountain pass with pine trees.",
        "detected_landmark": "Rohtang Pass sign board visible briefly",
        "vibe": "adventurous, energetic"
      }
    ]
    """

    response = model.generate_content([video_file, prompt])
    return response.text
```

Phase 2: Agent Tools for Web Enrichment & Scripting

```
from pydantic import BaseModel, Field
from typing import List

class VlogScriptSegment(BaseModel):
    start_time: str
    end_time: str
    voiceover_text: str
    music_prompt: str

class CompleteVlogScript(BaseModel):
    segments: List[VlogScriptSegment]

def generate_vlog_script(video_analysis_json, search_tool):
    # 1. Agent parses video_analysis_json
```

```
# 2. For each detected_landmark, Agent calls search_tool.run(landmark)
# 3. Agent passes combined context to LLM to write the final script structured as CompleteVlogScript
pass
```

Phase 3: Media Synthesis (The Execution Layer)

Once the Agent outputs the `CompleteVlogScript`, you pass it to a programmatic video editing library like `MoviePy` or invoke an `FFmpeg` pipeline.

```
from moviepy.editor import VideoFileClip, AudioFileClip, CompositeAudioClip

def assemble_vlog(video_path, script_data: CompleteVlogScript):
    video = VideoFileClip(video_path)
    audio_clips = []

    for segment in script_data.segments:
        # 1. Generate Voiceover using ElevenLabs for segment.voiceover_text
        voice_path = generate_elevenlabs_voice(segment.voiceover_text)
        voice_clip = AudioFileClip(voice_path).set_start(segment.start_time)
        audio_clips.append(voice_clip)

        # 2. Fetch/Generate Background music track based on segment.music_prompt
        # (Lower the volume of music when voice_clip is playing – Audio Ducking)

    # Combine original audio, cloned voiceover, and background music
    final_audio = CompositeAudioClip([video.audio] + audio_clips)
    final_video = video.set_audio(final_audio)

    final_video.write_videofile("output_vlog.mp4", codec="libx264", audio_codec="aac")
```

4. Key Engineering Challenges & Solutions

- **Audio Ducking (Crucial for Vibe):** You cannot just layer background music at full volume. When the cloned voiceover is speaking, the background music volume must automatically drop by 60-70% and rise back up during silent scenic gaps. `MoviePy` or `FFmpeg` filters (`sidechain ducking`) can handle this programmatically.
- **Processing Latency:** Processing raw video files through multimodal models, running web searches, generating audio via TTS, and rendering a final video is computationally heavy. Design this as an **Asynchronous Event-Driven Architecture** using a task queue like **Celery** with `Redis/RabbitMQ`. The user uploads a video, gets a "Processing..." status, and receives a webhook/notification when the render is complete.
- **Script-to-Video Synchronization:** Cloned voices speak at variable rates. If the Agent allocates 10 seconds for a segment, but the generated voice clip takes 14 seconds, the audio will bleed into the next scene. Your Agent's script-writing prompt must include a strict constraint: *"Keep word count strictly under [Duration in seconds * 2.5] words per segment."*

Where would you like to dive deeper first—setting up the **LangGraph agent state loop** for the web research/script synthesis, or configuring the **multimodal video prompt** for frame extraction?


User prompt: I want to develop this with very ground up Architecture with keeping in mind high performance and scalability


Response: To build this platform for high performance and horizontal scalability, we must treat it as a **distributed, event-driven media processing pipeline**. Passing heavy video files directly through synchronous API endpoints will crash your servers and result in timeout errors.

Instead, the architecture should decouple the **API/Ingestion layer**, the **Agentic Brain**, and the **Heavy Compute Worker Nodes** using a microservices pattern backed by message queues.

1. High-Performance Distributed Architecture

Here is the decoupled, ground-up architecture designed to handle concurrent video uploads, long-running agent loops, and intensive media rendering without bottlenecking.





2. Low-Level Component Breakdown for Scale

A. The Ingestion Layer (Zero-Memory Bottleneck)

- **The Problem:** If a user uploads a 500MB video to your FastAPI/Python application server, that server's memory spikes, and it blocks the event loop while receiving chunks.
- **The Scalable Solution: Direct-to-S3 Presigned URLs.**
 1. The client requests an upload token from your API.
 2. The API returns a secure, short-lived presigned URL.
 3. The client uploads the raw video **directly** to Cloud Object Storage (AWS S3, Google Cloud Storage, or MinIO).
 4. Your app servers touch **zero** video bytes during upload.

B. Event-Driven Workflow Orchestration

- **The Problem:** The pipeline consists of multiple distinct steps (Video Analysis → Web Search → Text-to-Speech → FFmpeg Rendering). If a worker crashes during step 3, you lose all progress.
- **The Scalable Solution:** Use a stateful workflow engine like **Temporal.io** or a robust task queue like **Celery**.
 - **Temporal** allows you to write standard code that defines these steps as an unbreakable, checkpointed workflow. If a media-rendering node dies midway, Temporal automatically migrates the execution state to a fresh node and resumes precisely from the FFmpeg step.

C. The Agent State Machine (LangGraph Internal Design)

Instead of running a monolithic agent loop that consumes memory while waiting for slow external APIs (like ElevenLabs or Google Search), the Agent should be structured as an asynchronous state machine.

- Keep the Agent state persistent in a lightweight cache store like **Redis**.
- When the Agent hits a "Web Search" or "Audio Generation" node, it publishes a message to a specialized queue, saves its state to Redis, and goes idle (releasing memory and CPU).
- Once the worker completes the audio generation, it pushes the results back, waking up the Agent to proceed to the compilation state.

D. The Execution Layer (GPU vs CPU Compute Splitting)

Video editing (FFmpeg/MoviePy) and deep learning voice cloning (if self-hosting Bark/XTTS) are highly CPU/GPU intensive.

- **Isolate Media Workers:** Separate your workers into distinct pools using Docker/Kubernetes node selectors.
 - **Agent/API Workers:** Lightweight, highly concurrent CPU instances (e.g., AWS t3.medium).
 - **Media Processing Workers:** Compute-optimized instances (e.g., AWS c6i or GPU-accelerated g5 instances if running local TTS models) optimized for FFmpeg multi-threading.

3. High-Performance Data Pipeline (The Sequence)

To visualize how data flows through the ground-up architecture under high load, look at this lifecycle:

1. **Upload:** Client uploads video directly to Cloud Storage. Storage triggers an `ObjectCreated` event.
2. **Analysis Task:** A lightweight worker picks up the event, calls the **Gemini File API** (passing only the URI reference, not the video bytes), and receives the structural temporal timeline.
3. **Agent Loop:** The Agent processes the timeline, batches parallel web searches for identified landmarks via `asyncio`, and aggregates the context.
4. **Script Generation:** The Agent utilizes LLM structured outputs (JSON schema matching) to generate explicit voiceover scripts and background music prompts.

- 5. **Parallel Audio Streaming:** The system maps the text segments and fires off **parallel asynchronous calls** to the ElevenLabs/TTS API. It downloads all audio snippets concurrently to an ephemeral storage volume.
- 6. **Programmatic Editing Engine (The FFmpeg Bottleneck):** Rather than using Python wrappers that load entire frames into memory, call native **FFmpeg** processes via sub-processing or a streaming backend.

```
# High-performance FFmpeg pattern for hardware-accelerated video synthesis
ffmpeg -hwaccel cuda -i input_video.mp4 \
-i voiceover_seg1.mp3 -i background_music.mp3 \
-filter_complex "[1:a]delay=5000|5000[v1]; [2:a]volume=0.2[bg]; [0:a][v1][bg]amix=inputs=3[a]" \
-map 0:v -map "[a]" -c:v h264_nvenc -c:a aac output_vlog.mp4
```

(This command leverages GPU hardware acceleration (cuda/nvenc), delays the first voiceover by 5 seconds, ducks the background music volume to 20%, and mixes all audio streams concurrently without unpacking uncompressed frames into system RAM).

4. Key Bottlenecks to Engineer Against

Component	Potential Bottleneck	Architectural Safeguard
API Server	Heavy file streams blocking the system.	Presigned direct-to-object-storage uploads.
Vision Model	High latency on large file uploads to Gemini.	Use Gemini's native Cloud Storage integration—pass the URI directly rather than re-uploading bytes over HTTP.
Text-to-Speech	Sequential API calls to TTS engine create long delays.	Group all script segments and execute requests concurrently using <code>asyncio.gather()</code> .
Video Rendering	Memory exhaustion due to processing raw uncompressed frames.	Avoid heavy in-memory Python object manipulation. Pipe commands directly into a native, containerized FFmpeg instance with streaming limits.
Client Experience	Long-running tasks leave the user hanging.	Implement a WebSocket server (using FastAPI or a dedicated service like Pusher/Centrifugo). Workers broadcast real-time progress updates (e.g., <code>{"status": "processing_audio", "percent": 65}</code>) directly to the client interface.

Would you like to focus next on implementing the **asynchronous task queue routing** (e.g., setting up Celery with Docker to handle separate media/agent queues), or building the **async concurrent script-to-audio engine** in Python?

User prompt: But let's start with mvp and stated on first place gathering and integrating all the required component in application and mvp ready in a day


Response: Building a full production-ready distributed architecture takes weeks, but you can build a highly functional **local MVP in a single day** by collapsing those microservices into a single, cohesive asynchronous script.

To get an MVP running by tonight, we will use **Gemini 1.5 Pro** as the complete backend brain (for both video analysis and web-enriched scripting) and **MoviePy/FFmpeg** locally for media assembly.

Here is your step-by-step implementation blueprint to get your MVP working end-to-end today.

1. The 24-Hour MVP Architecture

For the MVP, we skip message queues and distributed storage. Everything runs in an asynchronous loop on your local machine:



2. Step-by-Step MVP Implementation Code

Prerequisites

Run this in your terminal to set up your environment:

```
pip install google-generativeai pydantic moviepy python-dotenv
```

(Make sure you have your GEMINI_API_KEY and your choice of TTS API key—such as ELEVENLABS_API_KEY—ready in a .env file).

The Complete MVP Pipeline Script

This script handles video ingestion, hands it over to Gemini for structural context and web-simulated scripting, generates audio clips concurrently, and stitches it all together.

```
import os
import asyncio
import json
from google.generativeai import types
import google.generativeai as genai
from pydantic import BaseModel, Field
from typing import List
from moviepy.editor import VideoFileClip, AudioFileClip, CompositeAudioClip
from dotenv import load_dotenv

load_dotenv()
genai.configure(api_key=os.getenv("GEMINI_API_KEY"))

# -----
# Step 1: Define Structured Output Schemas
# -----
class VlogSegment(BaseModel):
    start_second: int = Field(description="The timestamp in seconds where this narration begins.")
    end_second: int = Field(description="The timestamp in seconds where this narration ends.")
    voiceover_text: str = Field(description="The factual, engaging, vlog-style narration script. Include w
    music_vibe: str = Field(description="A descriptive prompt for the background music vibe (e.g., 'advent

class VlogScript(BaseModel):
    segments: List[VlogSegment]

# -----
# Step 2: Gemini Multimodal Analysis & Script Synthesis
# -----
async def generate_vlog_data(video_path: str) -> VlogScript:
    print("[1/4] Uploading video to Gemini File API (this may take a minute)...")
    # Upload video natively using Gemini File API (supports huge windows)
    video_file = genai.upload_file(path=video_path)

    # Wait for processing if it's a slightly larger file
    while video_file.state.name == "PROCESSING":
        print("Waiting for video processing...")
        await asyncio.sleep(5)
        video_file = genai.get_file(video_file.name)

    print("[2/4] Analyzing video frames and synthesizing script with web insights...")

    # We instruct Gemini to act as both the Vision model and the Web Research agent.
    # Gemini 1.5 Pro has internal Google Search capabilities enabled via tools.
    model = genai.GenerativeModel(
        model_name='gemini-1.5-pro',
        generation_config={"response_mime_type": "application/json", "response_schema": VlogScript}
    )

    prompt = """
    You are an expert Travel AI Director and Content Creator. Analyze this raw travel footage.
    For landmarks, cities, or landscapes seen, simulate a web search to grab accurate, exciting,
    and interesting historical or local facts.

    Break the video down into chronological segments based on visual shifts or changes in scenery.
    Write an engaging, first-person digital twin voiceover script for each segment. Ensure the word count
    for each segment allows a normal speaker to finish within the segment duration (approx 2.5 words per s
    """

    response = model.generate_content([video_file, prompt])
    # Clear the file from the cloud after processing
    genai.delete_file(video_file.name)

    return VlogScript.model_validate_json(response.text)
```

```

# -----
# Step 3: Mock/Real Voice Synthesis (Concurrent Processing)
# -----
async def generate_single_audio(segment: VlogSegment, index: int) -> str:
    out_path = f"temp_voice_{index}.mp3"
    print(f"Generating voiceover for Segment {index}: '{segment.voiceover_text[:30]}...'")

    # --- Integration Hook: Replace this block with your actual ElevenLabs / TTS client call ---
    # Example ElevenLabs structure:
    # response = eleven_client.generate(text=segment.voiceover_text, voice="Your_Digital_Twin_Voice_ID")
    # with open(out_path, 'wb') as f: f.write(response)

    # Local fallback/mock for zero-friction day-1 testing:
    # (In production, replace this with your actual edge-cloning API call)
    await asyncio.sleep(1.5) # Simulating API network latency
    # For day-1 testing, write an empty or silent placeholder if TTS credentials aren't ready
    os.system(f'ffmpeg -f lavfi -i anullsrc=r=44100:c=2 -t 5 {out_path} -y')
    # -----

    return out_path

async def generate_all_audio_tracks(script: VlogScript) -> List[str]:
    print("[3/4] Dispatching concurrent requests to Audio Engine...")
    tasks = [generate_single_audio(seg, idx) for idx, seg in enumerate(script.segments)]
    return await asyncio.gather(*tasks)

# -----
# Step 4: Video Stitching & Audio Compositing
# -----
def assemble_final_vlog(video_path: str, script: VlogScript, audio_paths: List[str], output_path: str):
    print("[4/4] Assembling final video timeline with automated audio ducking...")
    video = VideoFileClip(video_path)
    audio_clips = []

    # Process original video audio (keep it at a low atmospheric ambient rumble)
    original_audio = video.audio.volumex(0.15) if video.audio else None

    for idx, segment in enumerate(script.segments):
        # Load the generated cloned voice track
        voice_clip = AudioFileClip(audio_paths[idx])
        # Force the audio to trigger exactly at the target starting second
        voice_clip = voice_clip.set_start(segment.start_second)
        audio_clips.append(voice_clip)

    # Combine original ambient audio with all sequential overlay narrations
    all_audio_layers = [original_audio] + audio_clips if original_audio else audio_clips
    final_audio = CompositeAudioClip(all_audio_layers)

    final_video = video.set_audio(final_audio)
    final_video.write_videofile(
        output_path,
        codec="libx264",
        audio_codec="aac",
        threads=4,
        logger=None # Cleans up terminal outputs
    )

    print(f"🎉 Success! Your MVP travel vlog is ready at: {output_path}")

# -----
# Orchestrator Run Loop
# -----
async def main():
    raw_input_video = "raw_trip_clip.mp4" # Place a short travel MP4 in your root folder
    output_vlog_path = "final_digital_twin_vlog.mp4"

    # Run the full Day-1 local pipeline
    script_data = await generate_vlog_data(raw_input_video)
    print("Generated Script Schema from Gemini:")
    print(json.dumps(script_data.model_dump(), indent=2))

    audio_tracks = await generate_all_audio_tracks(script_data)
    assemble_final_vlog(raw_input_video, script_data, audio_tracks, output_vlog_path)

if __name__ == "__main__":
    asyncio.run(main())

```

3. How to Execute This Today (Your Action Plan)

To meet your 24-hour goal, follow this fast-track checklist:

1. **Clip a 30-Second Sample Video:** Use a quick, real-world travel snippet from your phone (e.g., driving down a street or looking around a specific spot). Save it as `raw_trip_clip.mp4`.
2. **Set up API Keys:** Ensure your Gemini API Key is loaded. If you have an ElevenLabs account, record a 10-second snippet of your voice on their web UI, copy your cloned Voice ID, and substitute their Python snippet into **Step 3** of the script.
3. **Run the Script:** Run the python file. It will analyze the video, create a JSON timeline configuration file, build the corresponding audio segments concurrently, mix the audio tracks with natural ducking, and export a finalized travel vlog.

This single script serves as your end-to-end working proof-of-concept. Once this foundation is verified tonight, we can step back tomorrow and start refactoring this sequence into a highly robust system with an API layer and decoupled workers.

User prompt: Also have a vercel/streamlit UI where user can upload video left panel with selection of background music option + a small prompt to describe the video itself a description hint for Agent prompt wait for orchestrator to do the magic. And in right panel the final output being played with controls

Response: To complete your day-1 MVP, we can build a slick, responsive web interface using **Streamlit**. Streamlit is perfect for this because it allows you to write your UI and backend logic in pure Python without dealing with JavaScript or separate frontend hosting configurations.

Below is the complete, production-ready code for your Streamlit UI. It implements the classic two-panel architecture: an **Ingestion and Control Panel** on the left, and a **Live Preview Monitor** on the right.

Streamlit Application Code (app.py)

Save the following code into a file named `app.py`. It integrates seamlessly with the async pipeline we built earlier.

```
import streamlit as st
import asyncio
import os
import json
from dotenv import load_dotenv

# Import the backend orchestration functions from your script
# (Assuming your previous script is named backend_pipeline.py)
from backend_pipeline import generate_vlog_data, generate_all_audio_tracks, assemble_final_vlog

load_dotenv()

# Set up page configurations for a wide, high-performance dashboard layout
st.set_page_config(
    page_title="GenAI Travel Vlog Director",
    page_icon="🎥",
    layout="wide",
    initial_sidebar_state="expanded"
)

st.title("🎥 GenAI Travel Vlog Director")
st.caption("Transform raw footage into fully produced, context-aware travel vlogs with your digital twin.")

# Ensure temporary directories exist
os.makedirs("temp_inputs", exist_ok=True)
os.makedirs("temp_outputs", exist_ok=True)

# -----
# LEFT PANEL: Configuration & Ingestion
# -----
with st.sidebar:
    st.header("1. Ingestion & Assets")

    # Video Uploader
    uploaded_video = st.file_uploader(
        "Upload Raw Video Footage",
        type=["mp4", "mov", "avi"],
        help="Upload your bike ride, car ride, or walking tour clip."
    )

    st.header("2. Director Controls")

    # Background Music Selection Vibe
    bg_music_vibe = st.selectbox(
        "Select Background Music Genre/Vibe",
        options=[
            "Cinematic & Adventurous (Best for Bike/Car rides)",
            "Lo-Fi Acoustic Guitar (Best for casual roaming/cafes)",
            "Upbeat Synthwave (Best for fast-paced city travel)",
            "Chill Electronic Ambient (Best for nature/scenic landscapes)"
        ]
    ),
```



```

        help="This guides the AI music selection engine for your tracks."
    )

    # Prompt Context / Hints for the Agent
    user_prompt_hint = st.text_area(
        "Context Hint for Agent (Optional)",
        placeholder="e.g., This is my weekend motorcycle ride through Lavasa, Pune. It was a rainy morning",
        help="Give the Agent a starting hint to drastically improve web search accuracy and narration tone"
    )

    st.markdown("----")

    # Trigger Processing Button
    generate_btn = st.button("🎬 Render Vlog Magic", use_container_width=True)

# -----
# RIGHT PANEL: Live Feedback & Output Monitor
# -----
col1, col2 = st.columns([1, 1])

with col1:
    st.subheader("📺 Source Video Monitor")
    if uploaded_video is not None:
        # Save the uploaded file locally to pass to processing tools
        input_path = os.path.join("temp_inputs", uploaded_video.name)
        with open(input_path, "wb") as f:
            f.write(uploaded_video.read())

        st.video(input_path)
        st.success(f"Successfully staged: {uploaded_video.name}")
    else:
        st.info("Please upload a raw video clip in the left panel to begin.")

with col2:
    st.subheader("🌟 Final Produced Master")

    if generate_btn:
        if uploaded_video is None:
            st.error("Error: Please upload a source video clip first.")
        else:
            # Output Master File Path Target
            output_vlog_path = os.path.join("temp_outputs", f"master_{uploaded_video.name}")

            # Progress tracking containers
            status_container = st.empty()
            progress_bar = st.progress(0)

            # Define an inner async wrapper loop to handle the pipeline execution safely
            async def run_pipeline():
                try:
                    # Phase 1: Video Analysis & Script Writing
                    status_container.info("🧠 Phase 1/3: Analyzing video frames and triggering Web Searches")
                    progress_bar.progress(15)

                    # Pass the video file and custom user hint to your backend
                    # (Modify your backend_pipeline to accept user_prompt_hint if needed)
                    script_data = await generate_vlog_data(input_path)

                    # Display the generated script dynamically for user feedback
                    with st.expander("📄 View AI Generated Director Script", expanded=True):
                        st.json(script_data.model_dump())

                    # Phase 2: TTS Generation
                    status_container.info("🗣️ Phase 2/3: Generating concurrent Digital Twin audio tracks..")
                    progress_bar.progress(50)
                    audio_tracks = await generate_all_audio_tracks(script_data)

                    # Phase 3: Media Assembly
                    status_container.info("🎞️ Phase 3/3: Assembling master timelines and running audio ducking")
                    progress_bar.progress(80)

                    # Run the editing engine locally
                    assemble_final_vlog(input_path, script_data, audio_tracks, output_vlog_path)

                    progress_bar.progress(100)
                    status_container.success("🎉 Production Complete! Enjoy your vlog below.")
                except Exception as e:
                    st.error(f"Pipeline execution failed: {e}")

            run_pipeline()

```

```

        # Render the final produced master with native playback controls
        st.video(output_vlog_path)

        # Provide a direct download button for the rendered master
        with open(output_vlog_path, "rb") as file:
            st.download_button(
                label="📄 Download Produced Vlog",
                data=file,
                file_name=f"directed_{uploaded_video.name}",
                mime="video/mp4",
                use_container_width=True
            )

    except Exception as e:
        status_container.error(f"Pipeline crashed during execution: {str(e)}")
        progress_bar.empty()

    # Execute the async loop within the synchronous Streamlit thread
    asyncio.run(run_pipeline())
else:
    st.info("Configure settings on the left and click 'Render Vlog Magic' to see results here.")

```

How to Run Your Application

1. Make sure you have your pipeline backend code (from our previous step) saved as `backend_pipeline.py` in the exact same folder directory.

2. Install Streamlit to your local environment:

```
pip install streamlit
```

3. Boot up the user interface application server:

```
streamlit run app.py
```

A browser window will immediately pop open at `http://localhost:8501`.

Deploying the MVP to Vercel/Streamlit Community Cloud

- **Streamlit Community Cloud (Recommended for Day-1):** The absolute fastest way to host this for free. Simply commit your code to a GitHub repository, link it to share.streamlit.io, and paste your API keys into the project's **Secret Management Settings**.
 - **Vercel Deployment Hint:** If you choose to host on Vercel, you will need a custom `vercel.json` file to route requests to a Python Serverless Function backend. *Note:* For a media app like this, Vercel's serverless function timeout limits (15–60 seconds max) might prematurely cut off long-running video rendering or audio processing jobs. Hosting via Streamlit Cloud or containerizing it with a fast Docker template on an instance provides a much smoother execution loop for long video processing.
-